

Introducing Window Functions

PostgreSQL's documentation does an excellent job of [introducing the concept of Window Functions](#): a window function performs a calculation across a set of table rows that are somehow related to the current row. This is comparable to the type of calculation that can be done with an aggregate function. But unlike regular aggregate functions, use of a window function does not cause rows to become grouped into a single output row — the rows retain their separate identities. Behind the scenes, the window function is able to access more than just the current row of the query result.

Through introducing window functions, we have also introduced two statements that you may not be familiar with: **OVER** and **PARTITION BY**. These are key to window functions. Not every window function uses **PARTITION BY**; we can also use **ORDER BY** or no statement at all depending on the query we want to run. You will practice using these clauses in the upcoming quizzes. If you want more details right now, [this resource](#) from Pinal Dave is helpful.

*Note: You can't use window functions and standard aggregations in the same query. More specifically, you can't include window functions in a **GROUP BY** clause.*

Creating a Running Total Using Window Functions

Using Derek's previous video as an example, create another running total. This time, create a running total of `standard_amt_usd` (in the `orders` table) over order time with

no date truncation. Your final table should have two columns: one with the amount being added for each new row, and a second with the running total.

Creating a Running Total Using Window Functions

```
SELECT standard_amt_usd,  
  
       SUM(standard_amt_usd) OVER (ORDER BY occurred_at) AS running_total  
  
FROM orders
```

If you'd like another example of partitioning, check out the top answer from this Stack Overflow post: [“Partition By” Keyword](#)

Creating a *Partitioned* Running Total Using Window Functions

Now, modify your query from the previous quiz to include partitions. Still create a running total of `standard_amt_usd` (in the `orders` table) over order time, but this time, date truncate `occurred_at` by year and partition by that same year-truncated `occurred_at` variable. Your final table should have three columns: One with the amount being added for each row, one for the truncated date, and a final column with the running total within each year.

Ranking Total Paper Ordered by Account

Select the `id`, `account_id`, and `total` variable from the `orders` table, then create a column called `total_rank` that ranks this total amount of paper ordered (from highest to lowest) *for each account* using a partition. Your final table should have these four columns.

Ranking Total Paper Ordered by Account

```
SELECT id,  
  
       account_id,  
  
       total,  
  
       RANK() OVER (PARTITION BY account_id ORDER BY total DESC) AS total_rank  
  
FROM orders
```

Aggregates in Window Functions with and without ORDER BY

Run the query that Derek wrote in the previous video in the first SQL Explorer below. Keep the query results in mind; you'll be comparing them to the results of another query next.

```
SELECT id,  
  
       account_id,  
  
       standard_qty,  
  
       DATE_TRUNC('month', occurred_at) AS month,  
  
       DENSE_RANK() OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS dense_rank,
```

```
SUM(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS sum_std_qty,
```

```
COUNT(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS count_std_qty,
```

```
AVG(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS avg_std_qty,
```

```
MIN(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS min_std_qty,
```

```
MAX(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS max_std_qty
```

```
FROM orders
```

Now remove `ORDER BY DATE_TRUNC('month', occurred_at)` in each line of the query that contains it in the SQL Explorer below. Evaluate your new query, compare it to the results in the SQL Explorer above, and answer the subsequent quiz questions.

```
SELECT id,
```

```
account_id,
```

```
standard_qty,
```

```
DATE_TRUNC('month', occurred_at) AS month,
```

```
DENSE_RANK() OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS dense_rank,
```

```
SUM(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS sum_std_qty,
```

```
COUNT(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS count_std_qty,
```

```
AVG(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS avg_std_qty,
```

```
MIN(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS min_std_qty,
```

```
MAX(standard_qty) OVER (PARTITION BY account_id ORDER BY  
DATE_TRUNC('month', occurred_at)) AS max_std_qty
```

```
FROM orders
```

QUESTION 1 OF 3

What is the value of `dense_rank` in every row for the following `account_id` values?

QUESTION 2 OF 3

What is the value of `sum_std_qty` in the **first row** for the following `account_id` values?

Reflect

What is happening when you omit the `ORDER BY` clause when doing aggregates with window functions?

Use the results from the queries above to guide your thoughts then jot these thoughts down in a few sentences in the text box below.

Aggregates in Window Functions with and without ORDER BY

The `ORDER BY` clause is one of two clauses integral to window functions. The `ORDER` and `PARTITION` define what is referred to as the “window”—the ordered subset of data over which calculations are made. Removing `ORDER BY` just leaves an unordered partition; in our query's case, each column's value is simply an aggregation (e.g., sum, count, average, minimum, or maximum) of all the `standard_qty` values in its respective `account_id`.

As Stack Overflow user mathguy [explains](#):

The easiest way to think about this - leaving the `ORDER BY` out is equivalent to "ordering" in a way that all rows in the partition are "equal" to each other. Indeed, you can get the same effect by explicitly adding the `ORDER BY` clause like this: `ORDER BY 0` (or "order by" any constant expression), or even, more emphatically, `ORDER BY NULL`.

You can check out a complete list of window functions in Postgres (the syntax Mode uses) in the [Postgres documentation](#).

QUIZ QUESTION

Using Derek's example in the previous window, deconstruct the window function alias into its two parts: the *alias* part and the *window function* part.

Now, create and use an alias to shorten the following query (which is different than the one in Derek's previous video) that has multiple window functions. Name the alias `account_year_window`, which is more descriptive than `main_window` in the example above.

```
SELECT id,
```

```
    account_id,
```

```

DATE_TRUNC('year', occurred_at) AS year,

DENSE_RANK() OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS dense_rank,

total_amt_usd,

SUM(total_amt_usd) OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS sum_total_amt_usd,

COUNT(total_amt_usd) OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS count_total_amt_usd,

AVG(total_amt_usd) OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS avg_total_amt_usd,

MIN(total_amt_usd) OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS min_total_amt_usd,

MAX(total_amt_usd) OVER (PARTITION BY account_id ORDER BY
DATE_TRUNC('year', occurred_at)) AS max_total_amt_usd

FROM orders

```

Aliases for Multiple Window Functions

```

SELECT id,

account_id,

```

```
DATE_TRUNC('year', occurred_at) AS year,
```

```
DENSE_RANK() OVER account_year_window AS dense_rank,
```

```
total_amt_usd,
```

```
SUM(total_amt_usd) OVER account_year_window AS sum_total_amt_usd,
```

```
COUNT(total_amt_usd) OVER account_year_window AS count_total_amt_usd,
```

```
AVG(total_amt_usd) OVER account_year_window AS avg_total_amt_usd,
```

```
MIN(total_amt_usd) OVER account_year_window AS min_total_amt_usd,
```

```
MAX(total_amt_usd) OVER account_year_window AS max_total_amt_usd
```

```
FROM orders
```

```
WINDOW account_year_window AS (PARTITION BY account_id ORDER BY  
DATE_TRUNC('year', occurred_at))
```

Instructor Notes

Let's look at this again. We have broken down the syntax to explain LAG and LEAD functions separately.

LAG function

Purpose

It returns the value from a previous row to the current row in the table.

Step 1:

Let's first look at the **inner query** and see what this creates.

```
SELECT    account_id, SUM(standard_qty) AS standard_sum
FROM      orders
GROUP BY  1
```

What you see after running this SQL code:

1. The query sums the standard_qty amounts for each account_id to give the standard paper each account has purchased over all time. E.g., account_id 2951 has purchased 8181 units of standard paper.
2. Notice that the results are not ordered by account_id or standard_qty.

account_id	standard_sum
2951	8181
2651	4231
3401	116
2941	790
1501	7941
1351	3174
1721	1176
1191	6304

Step 2:

We start building the **outer query**, and name the inner query as `sub`.

```
SELECT account_id, standard_sum
```

```
FROM (
```

```
    SELECT account_id, SUM(standard_qty) AS standard_sum
```

```
    FROM orders
```

```
    GROUP BY 1
```

```
) sub
```

This still returns the same table you see above, which is also shown below.

account_id	standard_sum
2951	8181
2651	4231
3401	116
2941	790
1501	7941

Step 3 (Part A):

We add the Window Function `OVER (ORDER BY standard_sum)` in the outer query that will create a result set in ascending order based on the *standard_sum* column.

```
SELECT account_id,
```

```
       standard_sum,
```

```
       LAG(standard_sum) OVER (ORDER BY standard_sum) AS lag
```

```
FROM (
```

```
SELECT account_id, SUM(standard_qty) AS standard_sum
```

```
FROM orders
```

```
GROUP BY 1
```

```
) sub
```

This ordered column will set us up for the other part of the Window Function (see below).

Step 3 (Part B):

The LAG function creates a new column called *lag* as part of the **outer query**:

LAG(standard_sum) OVER (ORDER BY standard_sum) AS lag. This new column named *lag* uses the values from the ordered *standard_sum* (Part A within Step 3).

```
SELECT account_id,
```

```
standard_sum,
```

```
LAG(standard_sum) OVER (ORDER BY standard_sum) AS lag
```

```
FROM (
```

```
SELECT account_id,
```

```
SUM(standard_qty) AS standard_sum
```

```
FROM demo.orders
```

```
GROUP BY 1
```

```
) sub
```

Each row's value in *lag* is pulled from the previous row. E.g., for `account_id` 1901, the value in *lag* will come from the previous row. However, since there is no previous row to pull from, the value in *lag* for `account_id` 1901 will be NULL. For `account_id` 3371, the value in *lag* will be pulled from the previous row (i.e., `account_id` 1901), which will be 0. This goes on for each row in the table.

What you see after running this SQL code:

account_id	standard_sum	lag
1901	0	
3371	79	0
1961	102	79
3401	116	102
3741	117	116
4321	123	117
3941	148	123
1671	149	148

Step 4:

To compare the values between the rows, we need to use both columns (*standard_sum* and *lag*). We add a new column named *lag_difference*, which subtracts the *lag* value from the value in *standard_sum* for each row in the table:

```
standard_sum - LAG(standard_sum) OVER (ORDER BY standard_sum) AS
lag_difference
```

```
SELECT account_id,
```

```
standard_sum,
```

```
LAG(standard_sum) OVER (ORDER BY standard_sum) AS lag,
```

```

        standard_sum - LAG(standard_sum) OVER (ORDER BY standard_sum) AS
lag_difference

FROM (

    SELECT account_id,

        SUM(standard_qty) AS standard_sum

    FROM orders

    GROUP BY 1

) sub

```

Each value in *lag_difference* is comparing the row values between the 2 columns (*standard_sum* and *lag*). E.g., since the value for *lag* in the case of *account_id* 1901 is NULL, the value in *lag_difference* for *account_id* 1901 will be NULL. However, for *account_id* 3371, the value in *lag_difference* will compare the value 79 (*standard_sum* for *account_id* 3371) with 0 (*lag* for *account_id* 3371) resulting in 79. This goes on for each row in the table.

What you see after running this SQL code:

account_id	standard_sum	lag	lag_difference
1901	0		
3371	79	0	79
1961	102	79	23
3401	116	102	14
3741	117	116	1
4321	123	117	6
3941	148	123	25
1671	149	148	1

Now let's look at the **LEAD** function.

LEAD function

Purpose:

Return the value from the row following the current row in the table.

Step 1:

Let's first look at the **inner query** and see what this creates.


```
SELECT    account_id,  
  
          SUM(standard_qty) AS standard_sum  
  
FROM      demo.orders  
  
GROUP BY  1
```

What you see after running this SQL code:

1. The query sums the standard_qty amounts for each account_id to give the standard paper each account has purchased over all time. E.g., account_id 2951 has purchased 8181 units of standard paper.
2. Notice that the results are not ordered by account_id or standard_qty.

account_id	standard_sum
2951	8181
2651	4231
3401	116
2941	790
1501	7941
1351	3174
1721	1176
1191	6304

Step 2:

We start building the **outer query**, and name the inner query as `sub`.

```
SELECT account_id,
```

```
        standard_sum
```

```
FROM (
```

```
    SELECT account_id,
```

```
           SUM(standard_qty) AS standard_sum
```

```
    FROM demo.orders
```

```
    GROUP BY 1
```

```
) sub
```

This will produce the same table as above, but sets us up for the next part.

account_id	standard_sum
2951	8181
2651	4231
3401	116
2941	790
1501	7941
1351	3174
1721	1176
1191	6304

Step 3 (Part A):

We add the Window Function (`OVER BY standard_sum`) in the outer query that will create a result set ordered in ascending order of the *standard_sum* column.

```
SELECT account_id,
```

```
standard_sum,
```

```
LEAD(standard_sum) OVER (ORDER BY standard_sum) AS lead
```

```
FROM (

    SELECT account_id,

           SUM(standard_qty) AS standard_sum

    FROM demo.orders

    GROUP BY 1

) sub
```

This ordered column will set us up for the other part of the Window Function (see below).

Step 3 (Part B):

The `LEAD` function in the Window Function statement creates a new column called *lead* as part of the outer query: `LEAD(standard_sum) OVER (ORDER BY standard_sum) AS lead`

This new column named *lead* uses the values from *standard_sum* (in the ordered table from Step 3 (Part A)). Each row's value in *lead* is pulled from the row after it. E.g., for *account_id* 1901, the value in *lead* will come from the row following it (i.e., for *account_id* 3371). Since the value is 79, the value in *lead* for *account_id* 1901 will be 79. For *account_id* 3371, the

value in *lead* will be pulled from the following row (i.e., account_id 1961), which will be 102.

This goes on for each row in the table.

```
SELECT account_id,  
  
       standard_sum,  
  
       LEAD(standard_sum) OVER (ORDER BY standard_sum) AS lead  
  
FROM (  
  
    SELECT account_id,  
  
           SUM(standard_qty) AS standard_sum  
  
    FROM demo.orders  
  
    GROUP BY 1  
  
    ) sub
```

What you see after running this SQL code:

account_id	standard_sum	lead
1901	0	79
3371	79	102
1961	102	116
3401	116	117
3741	117	123
4321	123	148
3941	148	149
1671	149	183

Step 4: To compare the values between the rows, we need to use both columns (*standard_sum* and *lag*). We add a column named *lead_difference*, which subtracts the value in *standard_sum* from *lead* for each row in the table: `LEAD(standard_sum) OVER (ORDER BY standard_sum) - standard_sum AS lead_difference`

```
SELECT account_id,
```

```
    standard_sum,
```

```
    LEAD(standard_sum) OVER (ORDER BY standard_sum) AS lead,
```

```
    LEAD(standard_sum) OVER (ORDER BY standard_sum) - standard_sum AS
lead_difference
```

```
FROM (
```

```
SELECT account_id,
```

```
SUM(standard_qty) AS standard_sum
```

```
FROM orders
```

```
GROUP BY 1
```

```
) sub
```

Each value in *lead_difference* is comparing the row values between the 2 columns (*standard_sum* and *lead*). E.g., for *account_id* 1901, the value in *lead_difference* will compare the value 0 (*standard_sum* for *account_id* 1901) with 79 (*lead* for *account_id* 1901) resulting in 79. This goes on for each row in the table.

What you see after running this SQL code:

account_id	standard_sum	lead	lead_difference
1901	0	79	79
3371	79	102	23
1961	102	116	14
3401	116	117	1
3741	117	123	6
4321	123	148	25
3941	148	149	1
1671	149	183	34

Scenarios for using LAG and LEAD functions

You can use LAG and LEAD functions whenever you are trying to compare the values in adjacent rows or rows that are offset by a certain number.

Example 1: You have a sales dataset with the following data and need to compare how the market segments fare against each other on profits earned.

Market Segment	Profits earned by each market segment
A	\$550
B	\$500
C	\$670
D	\$730
E	\$982

Example 2: You have an inventory dataset with the following data and need to compare the number of days elapsed between each subsequent order placed for Item A.

Inventor y	Order_ i d	Dates when orders were placed
Item A	001	11/2/2017
Item A	002	11/5/2017
Item A	003	11/8/2017
Item A	004	11/15/2017
Item A	005	11/28/2017

As you can see, these are useful data analysis tools that you can use for more complex analysis!

Comparing a Row to Previous Row

In the previous video, Derek outlines how to compare a row to a previous or subsequent row. This technique can be useful when analyzing time-based events. Imagine you're an analyst at Parch & Posey and you want to determine how the current order's total revenue ("total" meaning from sales of all types of paper) compares to the next order's total revenue.

Modify Derek's query from the previous video in the SQL Explorer below to perform this analysis. You'll need to use `occurred_at` and `total_amt_usd` in the `orders` table along with `LEAD` to do so. In your query results, there should be four columns: `occurred_at`, `total_amt_usd`, `lead`, and `lead_difference`.

```
SELECT account_id,

        standard_sum,

        LAG(standard_sum) OVER (ORDER BY standard_sum) AS lag,

        LEAD(standard_sum) OVER (ORDER BY standard_sum) AS lead,

        standard_sum - LAG(standard_sum) OVER (ORDER BY standard_sum) AS
lag_difference,

        LEAD(standard_sum) OVER (ORDER BY standard_sum) - standard_sum AS
lead_difference

FROM (

SELECT account_id,

        SUM(standard_qty) AS standard_sum
```

```
FROM orders
```

```
GROUP BY 1
```

```
) sub
```

Comparing a Row to a Previous Row

```
SELECT occurred_at,
```

```
total_amt_usd,
```

```
LEAD(total_amt_usd) OVER (ORDER BY occurred_at) AS lead,
```

```
LEAD(total_amt_usd) OVER (ORDER BY occurred_at) - total_amt_usd AS  
lead_difference
```

```
FROM (
```

```
SELECT occurred_at,
```

```
SUM(total_amt_usd) AS total_amt_usd
```

```
FROM orders
```

```
GROUP BY 1
```

```
) sub
```

You can use window functions to identify what percentile (or quartile, or any other subdivision) a given row falls into. The syntax is `NTILE(*# of buckets*)`. In this case, `ORDER BY` determines which column to use to determine the quartiles (or whatever number of 'tiles you specify).

Expert Tip

In cases with relatively few rows in a window, the `NTILE` function doesn't calculate exactly as you might expect. For example, If you only had two records and you were measuring percentiles, you'd expect one record to define the 1st percentile, and the other record to define the 100th percentile. Using the `NTILE` function, what you'd actually see is one record in the 1st percentile, and one in the 2nd percentile.

In other words, when you use a `NTILE` function but the number of rows in the partition is less than the `NTILE(number of groups)`, then `NTILE` will divide the rows into as many groups as there are members (rows) in the set but then stop short of the requested number of groups. If you're working with very small windows, keep this in mind and consider using quartiles or similarly small bands.

You can check out a complete list of window functions in Postgres (the syntax Mode uses) in the [Postgres documentation](#).

[NEXT](#)

Percentiles with Partitions

You can use partitions with percentiles to determine the percentile of a specific subset of all rows. Imagine you're an analyst at Parch & Posey and you want to determine the largest orders (in terms of quantity) a specific customer has made to encourage them to

order more similarly sized large orders. You only want to consider the `NTILE` for that customer's `account_id`.

In the SQL Explorer below, write three queries (separately) that reflect each of the following:

1. Use the `NTILE` functionality to divide the accounts into 4 levels in terms of the amount of `standard_qty` for their orders. Your resulting table should have the `account_id`, the `occurred_at` time for each order, the total amount of `standard_qty` paper purchased, and one of four levels in a `standard_quartile` column.
2. Use the `NTILE` functionality to divide the accounts into two levels in terms of the amount of `gross_qty` for their orders. Your resulting table should have the `account_id`, the `occurred_at` time for each order, the total amount of `gross_qty` paper purchased, and one of two levels in a `gross_half` column.
3. Use the `NTILE` functionality to divide the orders for each account into 100 levels in terms of the amount of `total_amt_usd` for their orders. Your resulting table should have the `account_id`, the `occurred_at` time for each order, the total amount of `total_amt_usd` paper purchased, and one of 100 levels in a `total_percentile` column.

Note: To make it easier to interpret the results, order by the `account_id` in each of the queries.

Percentiles with Partitions

1.

SELECT

account_id,

occurred_at,

standard_qty,

NTILE(4) OVER (PARTITION BY account_id ORDER BY standard_qty) AS
standard_quartile

FROM orders

ORDER BY account_id DESC

2.

SELECT

account_id,

occurred_at,

```
gross_qty,
```

```
NTILE(2) OVER (PARTITION BY account_id ORDER BY gross_qty) AS gloss_half
```

```
FROM orders
```

```
ORDER BY account_id DESC
```

3.

```
SELECT
```

```
account_id,
```

```
occurred_at,
```

```
total_amt_usd,
```

```
NTILE(100) OVER (PARTITION BY account_id ORDER BY total_amt_usd) AS  
total_percentile
```

```
FROM orders
```

```
ORDER BY account_id DESC
```